

Topic Shift Detection for Chatbots: An Embedding-Centric Approach

1. Overview

We introduce a topics collection

Each topic stores its own summaries, embeddings, and message IDs.

When a new query arrives, we use vector similarity to determine:

- Should it belong to an existing research topic?
- Should a **new topic** be created?

This method is fast, deterministic, and scalable.

2. Database Schema

topics collection (new)

```
{  
  "_id": ObjectId,  
  "conv_id": ObjectId,  
  "topic_id": "topic_1",  
  
  "interaction_ids": [ObjectId, ...],  
  
  "long_summary": "Detailed summary of this topic's research-related  
  discussion.",  
  "short_summary": "1-2 line description of this research topic.",  
  
  "centroid_embedding": [0.23, -0.14, ...],  
  "entity_set": ["Google", "Q3 earnings", "Alphabet financials"],  
  
  "message_count": 12,  
  "created_at": ISODate,  
  "updated_at": ISODate}
```

}

Examples of entity sets:

- Topic “Google Earnings”
→ ["Google", "Alphabet", "Q3 revenue", "ad business", "cloud margins"]
 - Topic “BlackRock Macro Risk Outlook”
→ ["BlackRock", "risk commentary", "macro strategy", "bond yields"]
 - Topic “NVIDIA Client Note Q4”
→ ["NVIDIA", "GPU demand", "data center revenue", "client outlook"]
-

3. Embedding-Based Topic Classification

When a user interacts with the chatbot, each message is transformed into a **semantic embedding**. These embeddings naturally cluster according to the underlying research theme.

By representing each topic with a **centroid embedding**, we can:

1. Assign messages to existing clusters
2. Create new clusters when content diverges

This ensures topic purity even with highly interwoven research queries.

4. End-to-End Processing

4.1 Step 1 — Generate embedding for the new message

Example:

“What were BlackRock’s risk warnings for 2025 in their latest macro note?”

Compute normalized embedding:

$$e(m) \in \mathbb{R}^d$$

Where d = embedding dimension (e.g., 768).

4.2 Step 2 — Fetch all topics in the current conversation

These may include:

- Topic 1 → Google Earnings Discussion
 - Topic 2 → BlackRock Risk Outlook
 - Topic 3 → NVIDIA Client Quarterly Update
-

4.3 Step 3 — Compute similarity with each topic centroid

For each topic t with centroid μ_t :

Compute similarity

Interpretation:

- 0.80–0.95 → Same research domain

- 0.50–0.75 → Similar but not certain
- <0.50 → Likely unrelated

For example:

Topic	Centroid Similarity
Google Earnings	0.18
BlackRock Risk	0.91
NVIDIA Q4 Note	0.22

New query → assigned to “BlackRock Risk”.

4.4 Step 4 — Threshold Decision

Let t^* = topic with maximum similarity S^* .

If:

If $S^* > 0.75$, then

→ Assign message to existing topic t^* .

Else:

→ Create **new topic**, e.g.:

- “JP Morgan Equity Outlook”
 - “Macro Fed Interest Commentary”
 - “ESG Sector Probabilities”
 - “AI Equity Trend Analysis”
-

4.5 Step 5 — Update Topic Centroid

Given:

- μ_{\square} = old centroid
- $e(m)$ = new embedding
- n = previous message_count

Update via running average:

$$\mu_{t,\text{new}} = \frac{n \cdot \mu_{t,\text{old}} + e(m)}{n + 1}$$

Normalize:

$$\mu_{t,\text{new}} \leftarrow \frac{\mu_{t,\text{new}}}{\|\mu_{t,\text{new}}\|}$$

This keeps the centroid aligned with evolving discussions such as:

- Shifts in BlackRock rhetoric
 - Changes in NVIDIA guidance
 - Updates in Google revenue segmentation
-

4.6 Step 6 — Update Topic Document

Fields updated:

- `interaction_ids`
- `centroid_embedding`
- `entity_set`
- `long_summary`
- `short_summary`

- `updated_at`
 - increment `message_count`
-

5. Why This Works (Simple Mathematical Explanation)

1. Messages live in a “semantic space”

Each financial query:

- “Google’s cloud margins?”
- “BlackRock’s treasury yield assumptions?”
- “NVIDIA’s data center revenue?”

is represented by a vector $e(m)$.

2. Similar research questions form clusters

Queries about the same company/sector fall into dense groups.

Mathematically:

$$\text{Google Cluster} = \{e(m_1), e(m_2), \dots\}$$

$$\text{BlackRock Cluster} = \{e(m_k), e(m_{k+1}), \dots\}$$

3. The centroid represents topic meaning

$$\mu_t = \frac{1}{n} \sum_{i=1}^n e(m_i)$$

This is the “center of gravity” of the topic.

4. Cosine similarity identifies closest topic

$$\cos(e(m), \mu_t)$$

Highest similarity → best semantic match.

5. Threshold decides new vs existing topic

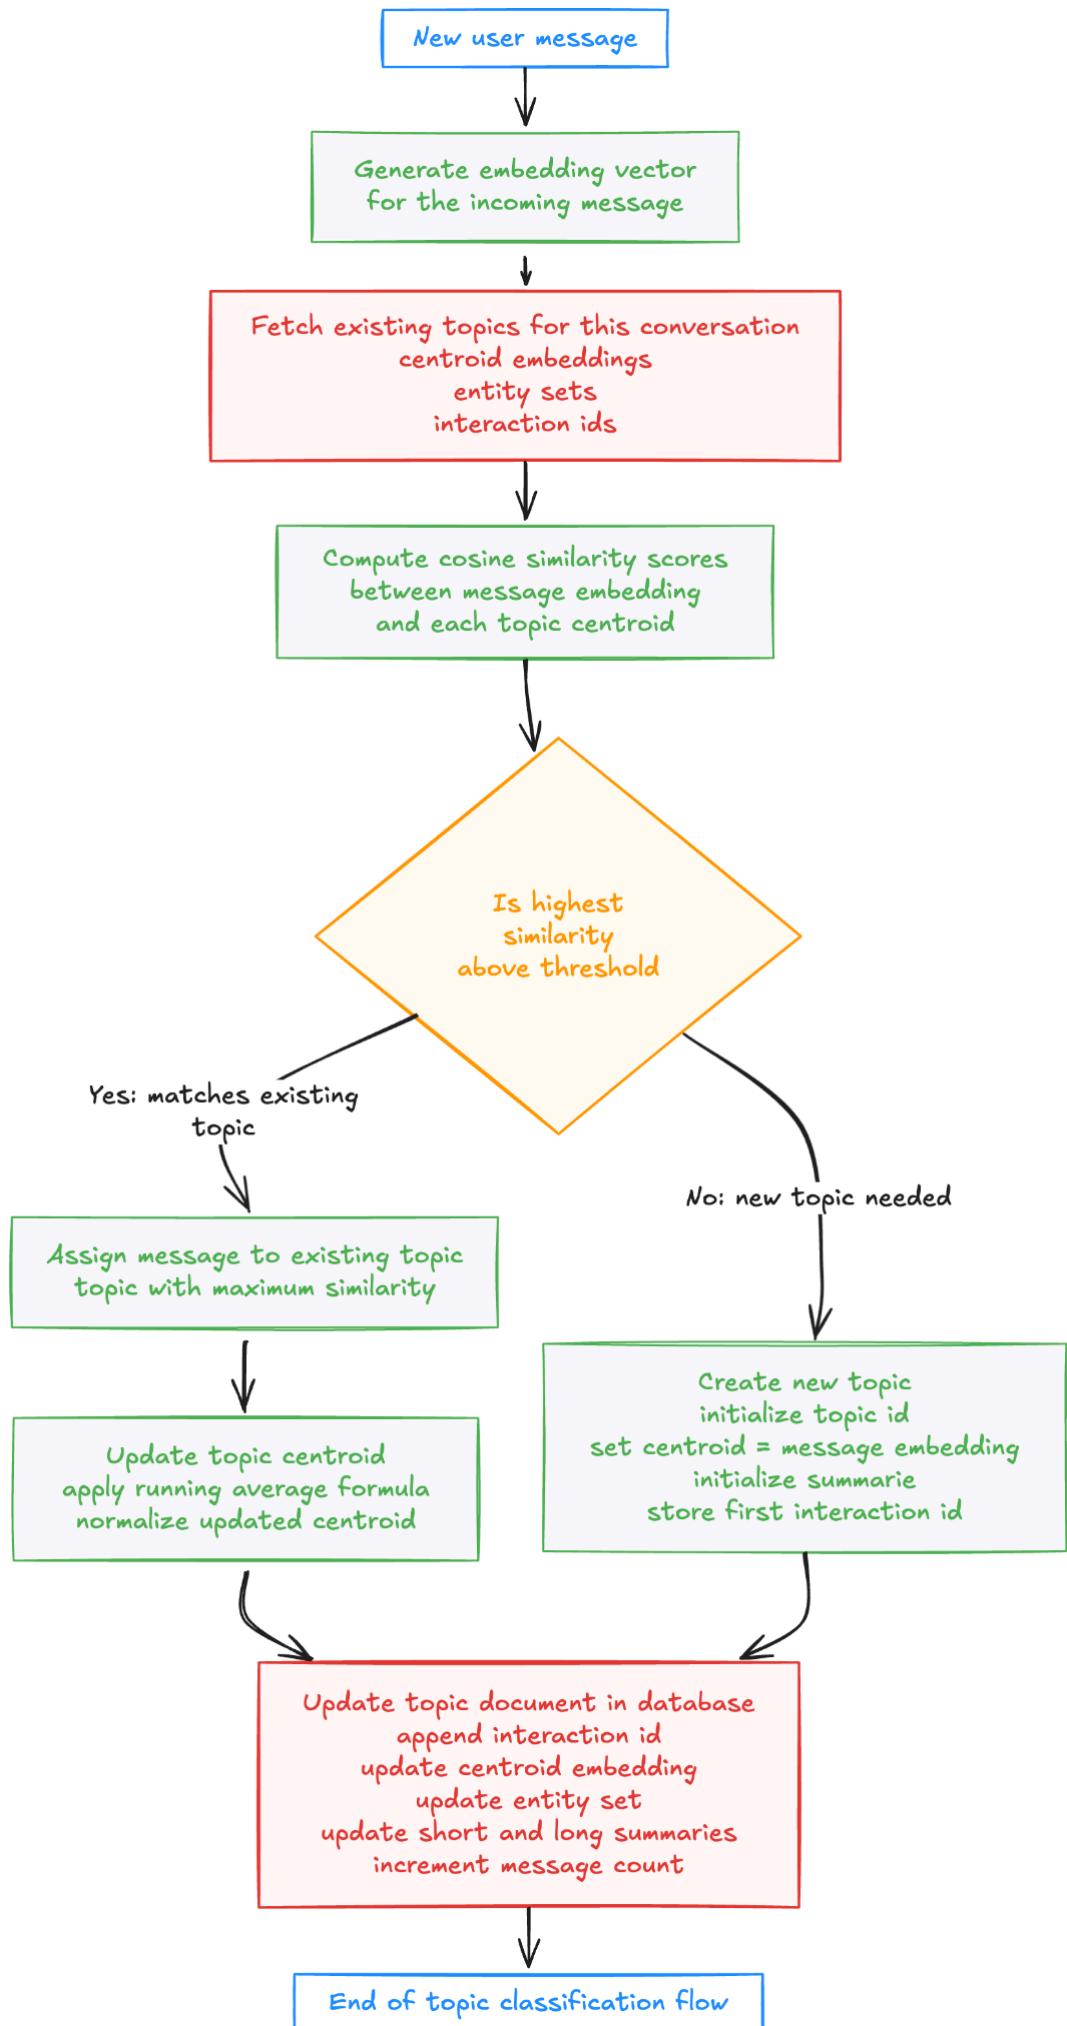
This avoids wrongly mapping:

- “Create a JP Morgan EM outlook brief”
to
- Google Earnings

Even if the conversation is interleaved like:

Google → BlackRock → Google → NVIDIA → JP Morgan

The vectors naturally separate these themes.



Simulating Dual-Process Thinking in Dialogue Topic Shift Detection

